

Sorting Report

Time: 11/03 – 16/03/2022

Student name: Nguyễn Duy Huy

Class: ATTN2021

Description

The purpose of this report is to measure and compare the execution time of Quicksort, Heapsort, Mergesort and `std::sort` of C++ STL.

Results

The first two test cases are increasing array and decreasing array. The others are random array.

Table of execution time (ms):

Test case	Quicksort	Heapsort	Mergesort	C++ STL Sort
0	25.866	45.755	19.35	7.899
1	25.142	52.578	25.59	5.419
2	71.166	139.09	81.071	50.646
3	71.37	145.876	79.032	49.77
4	71.967	143.968	79.136	50.568
5	72.524	134.464	80.088	50.895
6	72.539	146.555	80.112	50.489
7	72.328	140.056	79.401	50.185
8	72.419	143.849	79.554	50.125
9	72.066	141.267	79.352	50.92
Average	62.7387	123.3458	68.2686	41.6916

Chart of execution time:

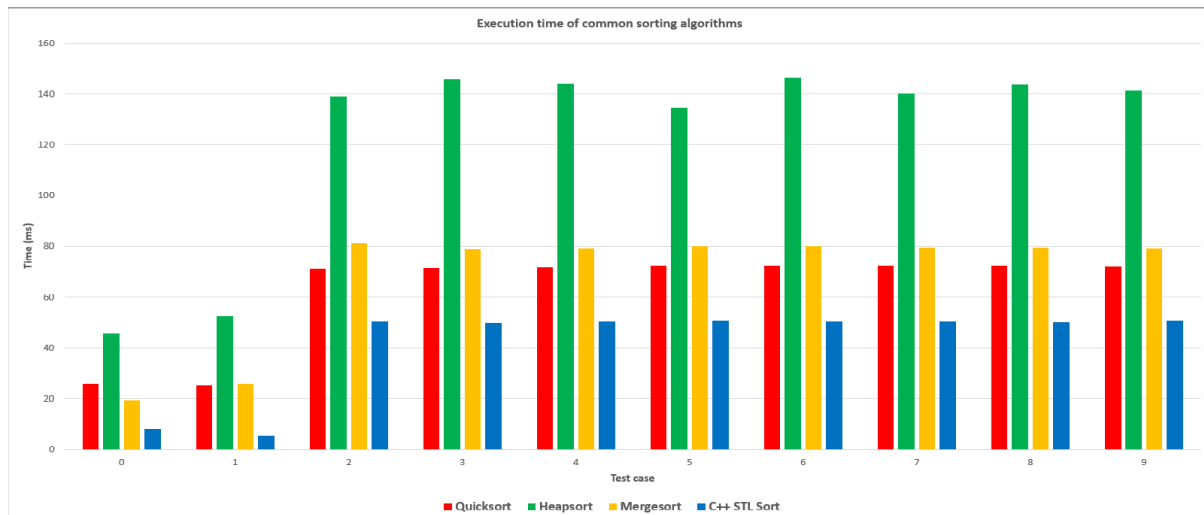


Table of instructions, cache-misses (less is better)

Algorithm	Instructions	Cache-misses
Quicksort	19717487159	1276707
Heapsort	22425146082	69918264
Mergesort	20597831400	2276122
C++ STL Sort	17847756639	881949

This benchmarking is run on:

- **OS:** Gentoo/Linux x86_64
- **Kernel:** Kernel: 5.15.23-gentoo-dist
- **CPU:** AMD Ryzen 5 5600H with Radeon Graphics (12) @ 3.300GHz
- **GPU:** AMD ATI Radeon RX 5500/5500M / Pro 5500M
- **GPU:** AMD ATI 07:00.0 Cezanne
- **Memory:** 8GB DDR4
- **g++ version:** (Gentoo 11.2.1_p20220115 p4) 11.2.1

According to those informations above we can easily realize that `std::sort` of C++ STL is far better than its opponents.

`std::sort`

It seems like they are all average $O(N \cdot \log(N))$ complexity, however, `std::sort` of C++ STL is better than the others.

Why `std::sort` is so quick? Because most of `std::sort` implementations use Quicksort, (or usually a hybrid algorithm like Introsort, which combines Quicksort,

Heapsort and Insertion sort).

So, `std::sort` is also Quicksort but better :))

Why Quicksort is faster than Mergesort?

Quicksort is faster than Mergesort because its cache performance is higher (less cache-misses) and it doesn't require extra space for merging operation as Mergesort.

Sometime, Quicksort has $O(N^2)$ in the worst case but we can avoid it by choosing random pivot.

How about Quicksort and Heapsort?

When both algorithms have same complexity ($\alpha * N * \log(N)$ for the Quicksort, and $\beta * N * \log(N)$ for the Heapsort), the Quicksort is faster because he has a proportionnality coefficient which equals the half of the Heapsort's proportionnality coefficient; mathematically, we have:

$$\alpha = \beta / 2$$

How to run the project?

Clone the project and build it:

```
$ git clone https://github.com/woanmeo11/sorting-benchmark.git
$ cd sorting-benchmark
$ ./build.sh
```

To display help, run the binary without any arguments:

```
$ ./benchmark
Usage: benchmark [OPTION]...
Measure execution time of common sorting algorithms.

--gentests      generate test cases

--all           equivalent to --heapsort --mergesort --quicksort --stl-sort
--heapsort      measure execution time of Heapsort
--mergesort     measure execution time of Mergesort
--quicksort     measure execution time of Quicksort
--stl-sort      measure execution time of std::sort of C++ STL
```

Then run `benchmark --gentests` to generate test cases:

```
$ ./benchmark --gentests
[*] Creating directory...
[*] Generating test case 0...
[*] Generating test case 1...
[*] Generating test case 2...
[*] Generating test case 3...
[*] Generating test case 4...
[*] Generating test case 5...
[*] Generating test case 6...
[*] Generating test case 7...
[*] Generating test case 8...
[*] Generating test case 9...
```

Measure the execution time

Run `benchmark --all` to view the benchmarking progress:

```
$ ./benchmark --all
```

Test case	Quicksort	Heapsort	Mergesort	C++ STL Sort
0	25.866	45.755	19.35	7.899
1	25.142	52.578	25.59	5.419
2	71.166	139.09	81.071	50.646
3	71.37	145.876	79.032	49.77
4	71.967	143.968	79.136	50.568
5	72.524	134.464	80.088	50.895
6	72.539	146.555	80.112	50.489
7	72.328	140.056	79.401	50.185
8	72.419	143.849	79.554	50.125
9	72.066	141.267	79.352	50.92

Show the number of instructions and cache-misses

To view the number of instructions and cache-misses, we use a tool called `perf`.
`--quicksort` option is used to investigate only Quicksort algorithm.

```
$ perf stat -e instructions,cache-misses ./benchmark --quicksort
```

Test case	Quicksort
0	24.092
1	24.47
2	71.317
3	71.76
4	71.92
5	72.016
6	72.353
7	72.016

7	72.912
8	72.681
9	71.411

Performance counter stats for './benchmark --quicksort':

19717487159	instructions:u
1276707	cache-misses:u

1.574716707 seconds time elapsed

1.553707000 seconds user

0.020004000 seconds sys

References:

- [heap-quick-comparison.pdf](#)
- [Memory location matters for performance](#)

Github:

- [sorting-benchmark](#)